



All seven 90-second PySR ablations · evaluation and generation comparisons

Mean across completed seeds ± 1 standard deviation

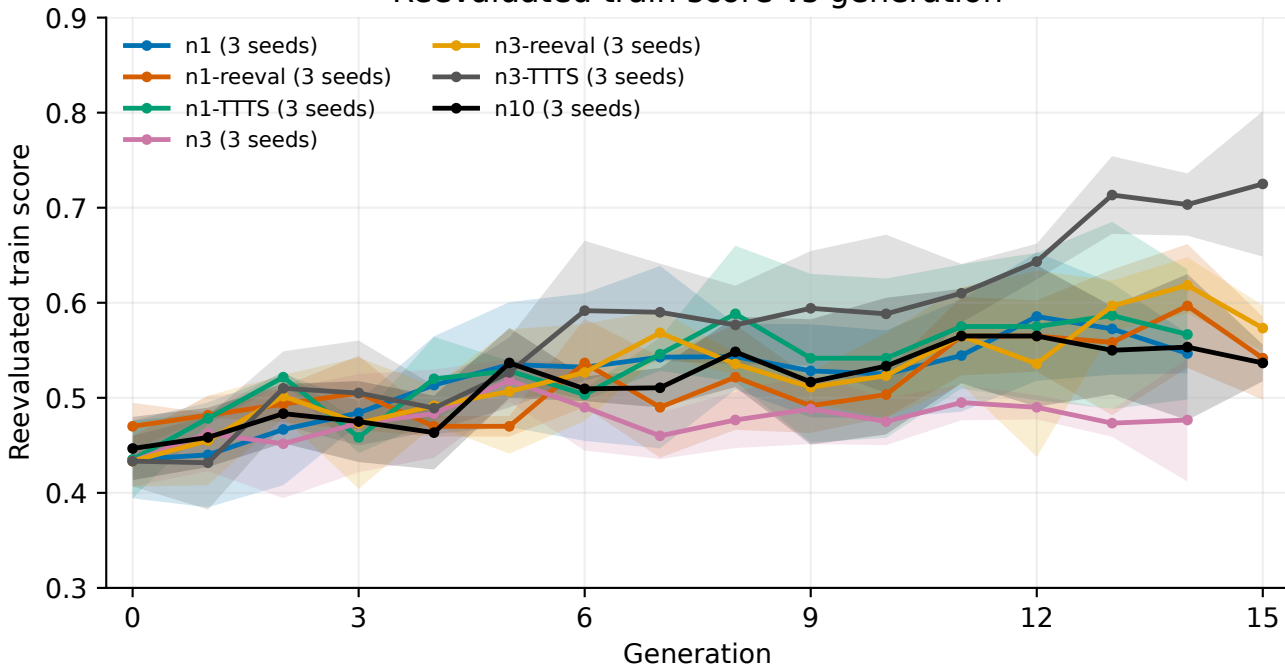
Reevaluated train score vs evaluations



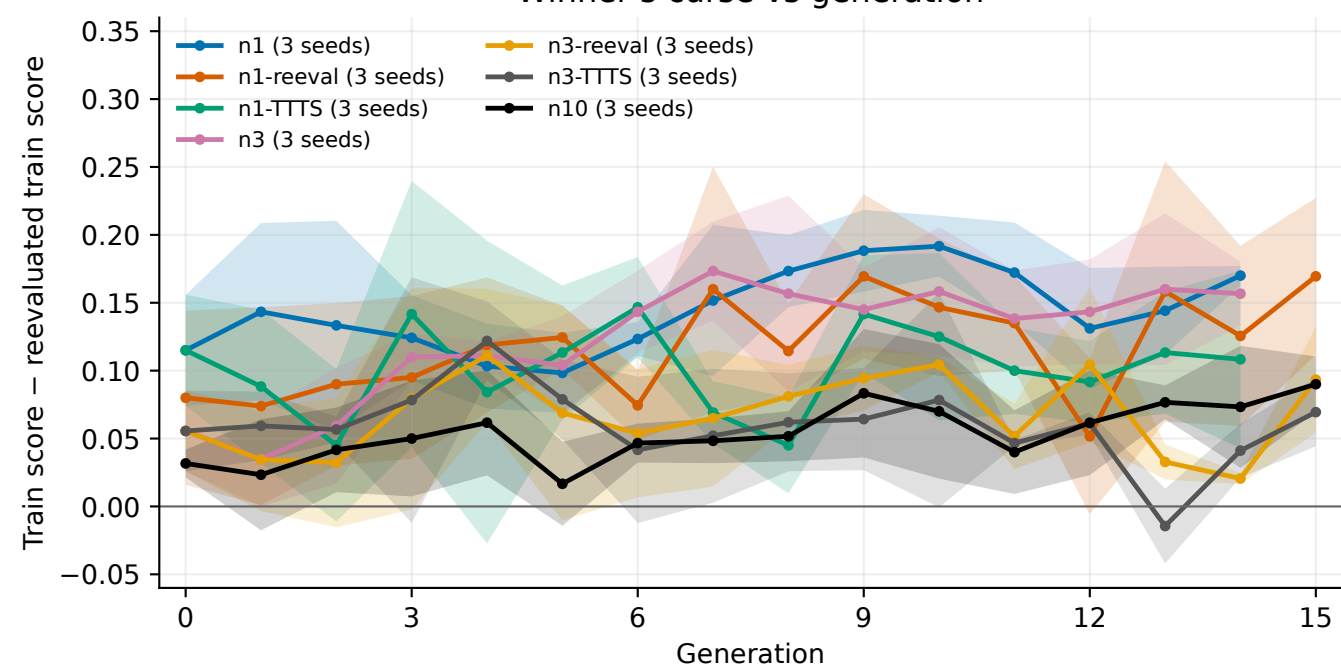
Winner's curse vs evaluations



Reevaluated train score vs generation



Winner's curse vs generation



Lines: seed mean; shading: ± 1 SD (ddof=0). One seed: no band. Linear interpolation within shared seed coverage; no extrapolation. Evaluations include selection reevaluations and exclude diagnostics. One seed-run covers the training task set.