



All seven 90-second PySR ablations · evaluation and generation comparisons

— Seed 1 - - - Seed 2


Reevaluated train score vs evaluations




Winner's curse vs evaluations



Reevaluated train score vs generation



Winner's curse vs generation



10 fresh seeds per best-candidate train diagnostic; missing diagnostics are not filled. Completed seeds only; resumed n3 seed 2 includes its earlier history. Evaluation counts include selection reevaluations and exclude diagnostics. One seed-run covers the training task set.