



90-second PySR ablations · best-candidate train reevaluation

Mean across completed seeds ± 1 standard deviation


n1: train reevaluation vs generation




n1: train reevaluation vs evaluations




n3: train reevaluation vs generation




n3: train reevaluation vs evaluations



n1: winner's curse



n3: winner's curse



Lines: seed mean; shading: ± 1 SD (ddof=0). One seed: no band. Linear interpolation within shared seed coverage; no extrapolation. Diagnostics use 10 fresh training seeds. Evaluations count evolution seed-runs across the training task set, excluding diagnostics.