

Table 1: SRBench results. Black-box R^2 is the mean over trials of the maximum test R^2 on each stored frontier (the –official convention). GT is the per-run solve percentage pooled over all datasets, seeds, and four noise levels. For the first three methods, black box uses R^2 -trained variants and GT uses GT-trained variants. — denotes unavailable or unreported results; MDLformer has no official entry. 1M denotes an evaluation cap, not a 1.5-minute time budget; wall-clock limits differ. PySR GT uses the complete original 1M run (290227); the current official selector instead picks an incomplete 15-minute restart portfolio. BasicSR has catastrophic negative- R^2 outliers; its mean is reported without clipping.

Metric	PySR++	BasicSR++	HPO	PySR	BasicSR	MDLformer
Black box R^2	0.821	0.819	0.818	0.798	-8.14×10^{10}	—
GT solve (%), 1M evals.	54.8	42.1	48.7	46.3	21.2	—
GT solve (%), 15 min	—	—	—	—	—	—
GT solve (%), 15 min + restarts	—	—	—	—	—	—